

AI for eGovernance: Combining Artificial Intelligence and Collective Intelligence to Develop Evidence-Based AI Policy

Dane Gambrell
The Governance Lab
Brooklyn, NY, United States
dane@thegovlab.org

Abstract—This case study describes how New Jersey's AI Task Force combined artificial intelligence (AI) with collective intelligence (CI) to develop evidence-based policy recommendations. The study discusses the role of AI in addressing two chronic challenges in policymaking: information overload and a lack of meaningful public engagement. By deploying Policy Synth, an open source deep research toolkit that uses AI agents to conduct automated web research and generate policy solutions, alongside the All Our Ideas online engagement tool to collect input from over 2,200 workers, the Task Force saved time and worked more inclusively. The findings suggest that AI-enabled research tools can enhance efficiency and the democratic legitimacy of policy outcomes. The Task Force's recommendations have already led to tangible policy actions, including a statewide AI workforce training initiative and the development of an AI-powered labor market monitoring system. The study helps to advance our understanding of the opportunities AI presents to transform policymaking by making the process faster and more efficient, enabling more meaningful engagement with communities, and widening the breadth of evidence used to inform policy, while also highlighting critical risks, such as algorithmic bias, transparency concerns, and the risk of overreliance on AI-generated insights. It also offers lessons for institutions who wish to integrate AI into public governance while underscoring the need for ethical and accountable AI deployment in policymaking.

Keywords—artificial intelligence (AI), collective intelligence (CI), reasoning models, AI agents, public engagement, policy development, digital democracy, civic technology, deep research

I. INTRODUCTION

In 2023, the State of New Jersey created an AI Task Force to help the state respond to the opportunities and challenges of artificial intelligence [1]. The Task Force's 21-member Workforce Training, Jobs of the Future, and Training Public Professionals Working Group ("the Working Group") was charged with examining the impact of AI on work and recommending actions to the Governor that the State and its companies, universities, unions, and other organizations can take to develop pathways to careers in the field of AI and responsibly upskill residents in AI.

The Working Group co-Chairs took an innovative approach to policymaking that combined artificial intelligence (AI)-enabled deep research tools with collective intelligence (CI) engagement methods to develop evidence-based policy recommendations for how the state can address AI's impact on

the state's workforce. This case study is written from the perspective of a policy professional who facilitated the research process and public engagement in support of the Working Group. It describes how the Working Group developed its recommendations and discusses the opportunities and challenges that AI-enabled research toolkits present for transforming the policymaking process.

II. THE CHALLENGE: OVERCOMING BARRIERS TO EFFECTIVE AI POLICYMAKING

The Working Group's use of AI and CI aimed to address two chronic barriers to effective and efficient policymaking: information overload and the lack of meaningful public participation. While these chronic challenges impact a range of policymaking areas, they are particularly acute for complex and rapidly changing problems like AI and its impacts on the workforce.

A. Information Overload: Finding Signal in the Noise

In today's information-rich environment, policymakers face a paradoxical challenge: drowning in data while struggling to gain meaningful insights that translate into effective policies. The amount of published research has grown rapidly in recent years, leaving researchers and scientists across fields overwhelmed [2]. While new knowledge is produced constantly, novel studies that could provide great value for policymakers are often buried beneath an avalanche of papers that merely reinforce existing viewpoints rather than offering fresh approaches and perspectives [3]. It is difficult for under-resourced policymakers to sift through the volume of information, distinguish quality research from speculation, reconcile contradictory findings, and translate abstract concepts into actionable policy recommendations.

B. Closed-Door Policymaking: The Need for Meaningful Public Engagement

Too often, the policymaking happens behind closed doors, with little meaningful impact from affected communities. One large-scale analysis of policy outcomes in the US political system concludes that "the preferences of the average American appear to have only a minuscule, near-zero, statistically non-significant impact upon public policy [4]." As a result, policies often reflect the interests and perspectives of the privileged rather than those of the broader public. This disconnect not only undermines the principles of democratic government but also

leads to policy solutions that fail to address the real-world concerns, needs, and experiences of diverse communities.

III. THE OPPORTUNITY: USING TECHNOLOGY TO REINVENT POLICYMAKING

At the same time, technological innovations have enabled new mechanisms for public participation in lawmaking [5]. Over the past two decades, lawmaking bodies around the world have experimented with using new information and communication technologies to facilitate a two-way conversation with the public [6]. A growing body of research within the field of collective intelligence (CI) – a field examining how groups work together to solve problems and make decisions – is dedicated to understanding how these experiments can be designed to improve the quality and legitimacy of lawmaking [7].

Since the introduction of powerful new large language models (LLMs), governments are exploring how AI can enhance the lawmaking and policymaking process. Brazil's Senate already uses technology to engage thousands of citizens in developing legislative proposal and is now investigating how AI could make citizens' participation more impactful [8]. California is deploying AI tools to gather ideas from residents about how to support communities impacted by the 2024 Los Angeles wildfires [9]. And the United Arab Emirates has announced ambitious plans to use AI to draft legislation and revise existing laws. Despite progress on incorporating technology into policymaking processes, such experiments are not yet widespread, and few efforts to institutionalize CI and AI experiments as part of public institutions' regular problem solving efforts have been successful [8]. While some studies have looked at how AI could enable new and better platforms for democratic deliberation, more experimentation is needed to understand how to effectively implement these processes in practice [10] [11].

IV. POLICY SYNTH AND ALL OUR IDEAS: TOOLS FOR SYNTHESIZING COLLECTIVE KNOWLEDGE

To address the twin challenges of information overload and a lack of meaningful participation opportunities for residents, the Working Group combined AI-enabled web research (sometimes called “deep research”) together with a large-scale public engagement to synthesize collective knowledge from both published research and insights from residents. This section describes two key technologies used by the Working Group: the deep research tool Policy Synth and the public engagement platform *All Our Ideas*.

A. About Policy Synth and Deep Research Toolkits

Policy Synth is a deep research toolkit that leverages artificial intelligence and genetic algorithms to make policy research more efficient, effective, and scalable [12]. Developed by *Citizens Foundation* and *The GovLab*, the toolkit is built as an open source JavaScript class-based library, meaning that its source code is published online under a license that allows for reuse and adaptation. The toolkit employs a human-centric approach that positions AI as a tool to augment human intelligence rather than replace it. (See [12] for a more detailed description of the technical design and architecture of Policy Synth.)

Deep research tools are AI-powered systems that “analyze vast amounts of data, synthesize information, and produce detailed outputs—including summaries, citations, and source links—in a fraction of the time it would take a human researcher [13].” These platforms go beyond simple search functionality by contextualizing information, identifying connections between sources, and generating comprehensive analyses that would traditionally require extensive human effort.

Deep research tools have been enabled by recent advancements in AI agents and reasoning models. AI agents are autonomous or semi-autonomous entities that interact with digital environments to accomplish tasks without the need for direct human intervention at each step. Reasoning models, such as OpenAI's o1 or DeepSeek's R1, are a type of LLM with enhanced abilities to follow multi-step reasoning processes, evaluate ideas against criteria, and generate nuanced analyses.

At its core, Policy Synth uses a collection of AI agents to automate and enhance the policy development workflow. First, its text compression agent takes lengthy text inputs and reduces them while maintaining essential details, supported by validation agents that check for correctness, completeness, and hallucinations. Second, it employs large-scale automated web research capabilities, using search engines to scan thousands of websites and analyze the data using both fast AI models and reasoning models. Third, Policy Synth implements Elo scoring to rank problems and solutions. Originally developed for chess player rankings, Elo scoring is an algorithmic ranking system where each problem or solution is compared in pairs, with rankings adjusted based on the outcome of each comparison [14]. This pairwise ranking functionality allows the system to avoid attempting to rank large sets all at once—a task that is challenging even for advanced LLMs.

Policy Synth also incorporates genetic algorithms to “evolve” solutions. The algorithm starts with an initial population of possible solutions that undergo selection based on a fitness function (using Elo ranking), followed by crossover (combining elements from different solutions) and mutation (introducing random changes). Through successive iterations, the algorithm refines raw solution concepts into more sophisticated policy proposals. At each stage the tool's ranking algorithm evaluates whether the changes improve the quality of the proposal – discarding those changes that do not. The evolutionary process is designed to emulate the ability of human experts to apply solutions from one domain to another, and to build upon one another's' ideas, as frequently happens during group deliberations [15].

Several commercial AI-powered research tools have emerged in recent years. These include the deep research platform Perplexity AI; Consensus, an AI-powered search engine tailored to scientific research; and Elicit, an AI platform that assists with developing literature reviews. OpenAI and Google have also integrated deep research functionalities into their respective ChatGPT and Gemini platforms.

What distinguishes Policy Synth from these commercial alternatives is: its origins as a tool designed specifically for policymaking, its open source development approach, customizability, and the integration of genetic algorithms to evolve solutions – features not commonly found in general-purpose deep research tools. The Working Group chose to develop Policy Synth as a custom solution tailored to its particular policy research needs.

The Working Group used Policy Synth to synthesize the collective wisdom contained in published sources about AI-driven changes to New Jersey’s workforce and economy into evidence-based policy recommendations.

B. About All Our Ideas and Digital Engagement Tools

All Our Ideas is an open source online engagement tool used to co-create a rank-ordered list based on public input. The platform has been deployed in thousands of consultations since its launch as a Princeton University research project in 2010, with applications ranging from urban planning initiatives to educational policy development. All Our Ideas is currently being renovated with a modern code base and retooled as a platform for social impact through a collaboration between Citizens Foundation and The GovLab.

All Our Ideas is a pairwise voting platform where individuals choose between a random pairing of binary choices in response to a prompt (See Fig. 1). Using the same Elo rating algorithm employed in Policy Synth, All Our Ideas dynamically ranks items based on how often they “win” their matchups against other ideas. Users can “vote” on pairings of statements as many or as few times as they wish. They can also skip a pairing if they do not wish to choose either statement. Users can also submit their own written response to the prompt; user-submitted responses can either be displayed as voting options for other users, or hidden according to the admin’s preference. The pairwise voting functionality allows organizations to efficiently collect data about preferences from large groups while also capturing qualitative data through user-submitted ideas. The platform generates a rank-ordered list in real time based on the results of the public voting (See Fig. 2).

While there are numerous digital platforms to facilitate online discussion, debate, and dialogue, [16], All Our Ideas specializes in preference aggregation and prioritization. This makes it particularly valuable for gathering opinions from residents during the problem prioritization phase of policy development, where the goal is to understand which issues matter most. It offers advantages in accessibility and inclusivity by lowering the participation threshold – users can contribute meaningful data with just a few clicks rather than writing lengthy responses. The quick, game-like interaction drives higher participation rates while still collecting substantive data. This combination of characteristics led the Working Group to identify All Our Ideas as the “best tool for the job.”

V. HOW IT WORKED: COMBINING PUBLIC ENGAGEMENT WITH DEEP RESEARCH

We undertook a four-stage process that combined public engagement with AI-enabled research to A) Identify problems, B) Prioritize among problems, C) Generate Solutions, and D) Transform solution ideas into implementable recommendations.

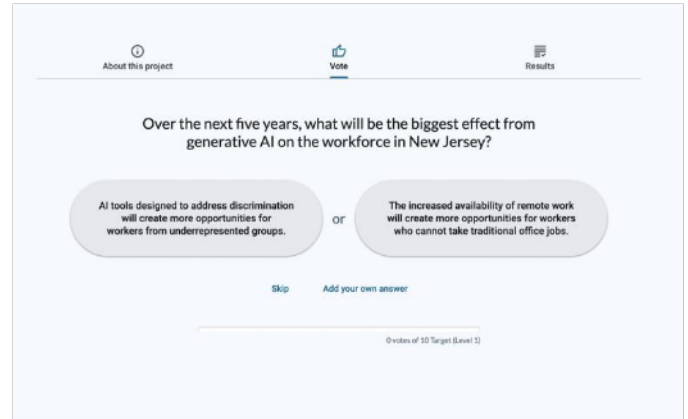


Fig. 1. Respondents selected between a random pairing of two ideas using the All Our Ideas platform.

A. Problem Identification

The Working Group began by defining the core challenge: understanding how AI might displace jobs in New Jersey’s workforce and economy. Policy experts in the Working Group first drafted and refined the problem statement – a short paragraph describing the challenge, why it is important, and its possible root causes.

Policy Synth’s AI agents then searched more than 1,000 online sources to identify potential impacts of AI on workers. AI agents crawled the Web for relevant information across various digital sources and generated a list of thousands of subproblems based on this search. The system then ranked these subproblems using Elo scoring to evaluate problems against criteria provided in custom instructions, including their economic impact, the number of people affected, and urgency within the next decade.

Through this ranking process, the AI agents transformed thousands of potential issues into a structured catalog of 15 problems, accompanied by references to the source materials from which each problem derived. The sub-problems covered areas such as the potential for AI to automate roles and displace jobs; change the skills that are in-demand; deepen inequality and economic disparities; and worsen existing biases in hiring and talent management. This AI-driven approach allowed the Working Group to search a wider range of sources than would be possible through traditional research methods, and to conduct this search and pull relevant findings more rapidly.

B. Problem Prioritization

The Working Group then engaged workers in New Jersey in refining its understanding of the problem using All Our Ideas. The tool presented respondents with two answer choices in response to the prompt: “Over the next five years, what will be the biggest effect from generative AI on the workforce in New Jersey?” Participants selected between random pairs of statements as many or as few times as they wished.

The list of 96 answer choices covered 59 challenges (such as job displacement and threats to worker power) and 37 opportunities (such as the potential of AI to create new job opportunities or to aid in job matching.) Over three-weeks, more than 2,200 private sector workers across New Jersey helped to rank these challenges [17]. The platform generated a real-time

rank-ordered list showing the most and least frequently selected options.

The prioritized list of 20 problem areas surfaced workers' concerns such as the potential for GenAI to displace jobs or lead to greater economic inequality. While the priority problem areas were similar to those identified by the AI-enabled web research, the engagement also surfaced additional challenges, such as concerns about the impacts of AI on older workers or privacy and surveillance concerns that were not ranked highly by the AI. (View the full rank-ordered list of problems [here](#)).

Combining outputs from AI-enabled research with direct public input created a more comprehensive understanding of the challenge, helping to guide the working group towards solutions that respond to real problems faced by workers.

C. Solution Generation

After prioritizing among problem areas, the Working Group used Policy Synth to develop solutions. We generated two sets of solutions in parallel to address 1) problems identified by the AI-enabled problem search and 2) problems identified by workers.

First, Policy Synth's AI agents developed solutions to the 15 highest-ranked problems identified through the AI-enabled web research. Similar to the problem search phase, AI agents crawled websites, academic journals, white papers, and other online sources, then wrote up the ideas using an LLM. The system then used its genetic algorithm to "evolve" these solutions through an iterative improvement process [18]. Through fifteen rounds of crossovers, mutations, and evaluations, the system refined solution concepts into 1,451 policy proposals covering 15 problem areas. Then, we repeated this process to generate solutions to the 20 highest-ranked problems as identified through the public engagement. This second search, solution generation and evolution process generate an additional 1,101 policy proposals covering the 20 highest-priority problem areas identified by workers.

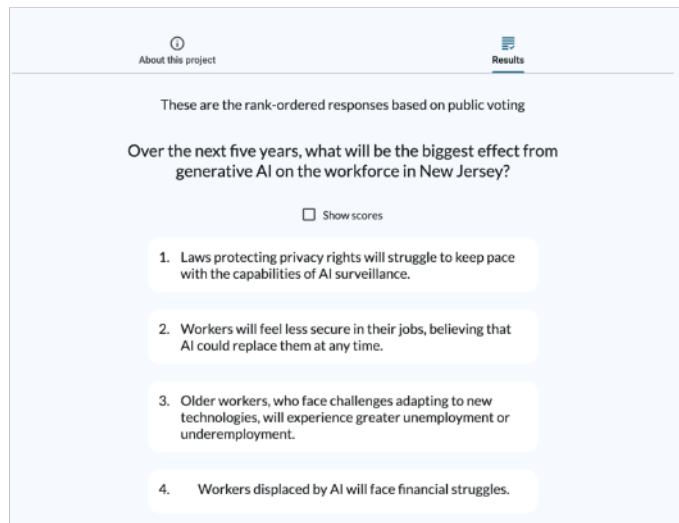


Fig. 2. Based on thousands of responses, All Our Ideas generated a rank-ordered list of the priority problem ideas.

The combined output from both searches was a catalog of more than 2,500 solution proposals covering all priority problem areas. Each solution was presented in a structured format that outlined its implementation strategy, expected outcomes, and potential challenges, providing the Working Group with comprehensive information for the final recommendation development phase.

D. Transforming Ideas Into Recommendations

Finally, we converted the solution proposals into actionable policy recommendations. The Working Group focused its recommendations on eight priority problem areas that were identified through the AI-enabled research and public engagement:

1. Economic Disparity
2. Reduction in Overall Labor Demand
3. AI-driven Occupational Shifts
4. Worker Displacement
5. Lack of Transparency and Accountability in AI Systems
6. Impacts on Mental Health
7. Impacts on Older Workers
8. Privacy and Surveillance Challenges.

For each of the eight priority areas, the Working Group selected the five highest-ranked solution proposals as determined by the prior Elo scoring. This prioritization narrowed the initial set of more than 2,500 proposals to a shortlist of 40 candidates for further consideration. Labor economists serving on the Task Force distilled the shortlisted solutions into four key recommendations for state leadership:

1. Expand AI-integrated Skill Development: Enhance education and workforce training programs to incorporate AI skills development opportunities for all New Jersey residents
2. Enhance NJ Career Navigator: Develop an AI-powered labor market monitoring and response system to help workers navigate changing career landscapes
3. Expand Transition Support: Strengthen state workforce programs that support workers through job transitions
4. Enable Small Business AI Adoption: Help Small Office/Home Office (SOHO) businesses leverage generative AI to drive growth.

Each recommendation was grounded in evidence from a literature review encompassing three dozen articles examining AI's impact on skills demand and worker transition strategies [19].

The recommendations have already resulted in concrete outcomes. In response to Recommendation 1, the state developed a free, comprehensive AI skills training program for all public servants in New Jersey [20]. As suggested in Recommendation 2, the state is developing an AI-powered labor

market monitoring system to help workers navigate career changes.

VI. BALANCING PROMISE AND PERIL: THE ROLE OF AI-ENABLED RESEARCH IN EVIDENCE-BASED POLICYMAKING

The Task Force's approach shows that AI-enabled research tools present both opportunities and challenges for policymaking. With the growing proliferation of commercial AI research platforms, it is increasingly important that we understand the benefits and risks posed by these technologies. This section explores these dual aspects and their implications for the future of evidence-based policy development.

A. Opportunities: Enhancing Policy Research Through AI-Human Collaboration

The use case demonstrates the potential of AI tools to transform policymaking by making the process faster and more efficient, enabling more meaningful engagement with communities, and widening the breadth of evidence used to inform policy.

First, the experience using Policy Synth demonstrated how AI can enable policymakers to conduct policy research more efficiently and quickly than traditional bench research processes would allow. We effectively used AI to process thousands of online sources and transformed the findings into implementable recommendations in just eight weeks with only one part-time policy professional supporting the Working Group. For policymakers facing increasingly complex challenges with limited resources, AI research tools offer a way to quickly develop a good understanding of multifaceted issues while maintaining evidence-based rigor.

Second, the case study demonstrates how AI can be used to synthesize inputs from web research together with findings from public engagement to create a more comprehensive understanding of the problem landscape. When the Working Group compared the AI-generated priority problems with those identified through public engagement, we discovered important complementary insights—the public highlighted concerns about older workers and privacy issues that weren't prioritized in the AI-generated analysis, while the AI research identified technical nuances that might not have emerged through public input alone. AI transformed what would typically be a manual, subjective process of reconciling different information sources into a structured, evidence-based approach to solution development.

Finally, the Task Force experience showcases the powerful synergy that can emerge when combining AI capabilities with human expertise in the policymaking process. Policy experts reviewed outputs at each stage and adapted the findings to New Jersey's specific workforce needs, economic landscape, and policy environment. This human-in-the-loop model shows how policymakers can use AI to expand the breadth of evidence used to inform policy while maintaining the irreplaceable roles of human creativity, reasoning, institutional knowledge, and decisionmaking ability.

B. Challenges: Addressing Bias, Opacity, Echo Chambers, and Overreliance on AI

A significant challenge for the adoption of AI-enabled deep research tools in policymaking is the potential for these systems

to perpetuate biases. AI agents may inadvertently reflect existing biases in data, literature, or other source materials, or biases within the algorithms that power them, in turn entrenching disparities rather than addressing them. AI systems used in policymaking must be designed, developed, and deployed according to ethical principles and best practices designed to mitigate bias and discrimination. Public institutions should also establish accountability mechanisms to evaluate AI-informed decision-making to ensure that outcomes are equitable and inclusive [21].

A related challenge is how to ensure transparency, explainability, and accountability when using AI tools to inform policy recommendations. Many commercial AI tools are “black box” systems, where users can understand inputs and outputs, but how the systems process data and arrive at decisions is a mystery [22]. As an open source toolkit, Policy Synth provides more transparency when compared to proprietary tools. At the same time, the complex nature of Policy Synth's genetic algorithms and ranking processes make it difficult to track and explain exactly how policy solutions were developed and prioritized. Public institutions may need to develop new frameworks for documenting AI's role in policy development, clearly delineating where human judgment overrides algorithmic recommendations, establishing responsibility boundaries when policies informed by AI systems fail to deliver expected results, and prioritize using AI systems that are explainable and interpretable [23].

As AI-enabled research tools become more sophisticated and widespread, there is a risk that powerful tools may be misused. One danger is that overtaxed and under-resourced policymakers may increasingly defer to algorithmic outputs rather than applying critical human judgment. U.S. federal agencies are increasingly using AI tools to inform a wide range of policy and personnel management decisions in areas ranging from healthcare [24] to immigration [25] to housing [26] – often without proper guardrails in place to prevent and rectify improper use. While AI tools are able to produce increasingly detailed, sophisticated, and authoritative-sounding reports and recommendations, observers have cautioned that these tools can still fabricate facts, generate fictitious citations, fail to convey its level of certainty [27], and cite misleading or poor-quality sources of information [28] – particularly when citing news content [29]. Early evaluations have suggested that reasoning models – the technology powering many AI search tools – may hallucinate at higher rates than older LLMs [30]. There is an urgent need for public institutions to establish clear protocols that position AI as an advisor rather than a decisionmaker and develop processes to align AI-generated recommendations with broader societal and ethical values.

Finally, deep research may create an echo chamber effect that prioritizes established beliefs over novel ideas. Search algorithms used to scan research databases often prioritize content with high visibility, engagement metrics, or established credibility markers [31]. Consequently, AI crawlers, which use search algorithms to collect information from digital sources, may over-sample from mainstream sources while failing to capture scholarship that is new or that presents alternative viewpoints, materials not readily accessible in digital formats, or knowledge published behind paywalls. Further, many website

are now blocking web crawlers to prevent AI developers from using sites' data to train their models [32]. These blindspots could impair innovation in policymaking by reinforcing conventional wisdom rather than facilitating the exploration of creative solutions. Policymakers seeking transformative and novel solutions may thus need to supplement AI-enabled research with intentional strategies to diversify their information sources.

C. Replication: Enabling Conditions and Context

Several conditions allowed the Working Group to successfully experiment with AI-enhanced policy development. Jurisdictions seeking to replicate this approach should consider how these conditions apply to their policymaking context.

Tech Infrastructure: First, in contrast to many public institutions which are resistant to new ways of working, New Jersey has made thoughtful investments in infrastructure and capacity building to enable the regular and repeated use of technology for public engagement. The state established an Office of Innovation in 2018 [33] to provide agencies across state government a dedicated team to support modernization and innovation initiatives, and appointed its first Chief AI Strategist in 2024 (who co-Chaired the Working Group and directed the AI-enabled research initiative) to provide agencies guidance on their implementation of AI. The Office of Innovation supported the Working Group's engagement with workers by acquiring and vetting (all tools used by the government need to be scrutinized for accessibility and security). The project also built on past experience. In 2020, the state's Future of Work Task Force used the All Our Ideas platform to engage New Jersey workers in prioritizing among problems created by the impact of emerging technologies on job skills, rights and benefits for workers, and workplace health and safety. This experience created familiarity with the platform, institutional knowledge, and trust in the approach among public servants, allowing the Working Group to replicate the engagement strategy with minimal barriers. By prioritizing continuous digital transformation and iterative improvement rather than investing in one-off solutions and technological silver bullets, New Jersey has developed the infrastructure, institutional knowledge, and innovation mindset needed to replicate the approach to address future challenges.

Leadership: A second condition that enabled the project to succeed was that it had champions in key decisionmaking positions. The state's Chief AI Strategist championed the innovative approach. But support from the Governor's Office and co-Chairs from the Department of Labor and the Office of the Secretary of Higher Education were equally crucial to ensuring rapid agreement to try this novel approach. Leadership and staff expertise also helped to ensure that outputs were reviewed for bias and relevance. The Working Group consulted decisionmakers early in the process to ensure that the final outputs were delivered in a form and format that the state could meaningfully act upon, ensuring process and its outcomes received serious consideration. Innovative projects too often fail because they lack an owner or champion within the institution who is committed to consider and use the results, because the ideas are not aligned to the needs and priorities of policymakers, or because insights are delivered in a level of detail or

abstraction that makes them inaccessible to policymakers [34]. Jurisdictions seeking to replicate this approach should identify and secure support from executive leadership who can advocate for the project and meaningfully use its results.

Lean Innovation: Finally, the Working Group chose tools that were well-suited to its priorities and needs. Policy Synth's genetic algorithms and evolutionary approach were particularly well-suited to the complex, multifaceted nature of AI workforce impacts, while All Our Ideas' pairwise voting methodology enabled efficient prioritization of issues by residents. Both tools are open source and highly customizable, allowing the Working Group to adapt the platforms to their unique needs. In addition, having staff trained in the use of these tools to conduct the deep research and organize the public engagement in house (in contrast to another Task Force Working Group that hired a large consultancy to conduct research at a significant cost) was essential. However, these tools are not a "one-size-fits-all" solution to any need. Policy Synth, in particular, requires a significant time and financial investment, as well as some technical skills, to use effectively. Jurisdictions may be able to achieve sufficient results with newly available deep research platforms. The key consideration should be matching technological capabilities to the needs of policymakers for evidence-based information [35].

VII. CONCLUSION

The New Jersey AI Task Force's innovative approach to AI workforce policy development suggests that combining AI-enabled research tools with collective intelligence methods can enhance policymaking for complex technological challenges. By leveraging AI to conduct large-scale automated research alongside large-scale online public engagement, the Working Group effectively addressed the twin challenges of information overload and meaningful public participation. This hybrid approach enabled the Task Force to rapidly develop evidence-based recommendations that are already yielding concrete results, including free AI skills training for all New Jersey public servants and the development of an AI-powered labor market monitoring system. The case illustrates how AI can serve as a powerful complement to human expertise in policy development, allowing for more efficient, comprehensive, and responsive governance.

However, policymakers must carefully consider the limitations and potential risks, including algorithmic bias, transparency, accountability, and overreliance on AI systems. Future research should explore how to design AI systems that better mitigate these risks while maximizing their benefits for public policy development. As governments increasingly adopt AI tools for policymaking, developing ethical principles, accountability mechanisms, and protocols that position AI as an advisor rather than a decision-maker will be essential for ensuring these technologies enhance rather than undermine democratic governance. The New Jersey case offers valuable lessons for other states and jurisdictions seeking to harness AI's potential while navigating its challenges in developing effective public policy.

REFERENCES

- [1] "Executive Order No. 346," The Office of New Jersey Governor Philip D. Murphy, October 2023, [Online] Available at: <https://d31hzhk6di2h5.cloudfront.net/20231010/e7/5b/7d/96/a898d65e01b6941bfd77839f/EO-346.pdf>
- [2] J. M. A. Hanson, P. G. Barreiro, P. Crosetto, D. Brockington, "The strain on scientific publishing," QSS, 1-21 (2024), [Online] Available at: <https://doi.org/10.48550/arXiv.2309.15884>
- [3] J. S. G. Chu and J. Evans, "Slowed canonical progress in large fields of science," PNAS vol. 118 no. 41, [Online] Available at: <https://doi.org/10.1073/pnas.2021636118>
- [4] M. Gilens and B. I. Page, "Testing Theories of American Politics: Elites, Interest Groups, and Average Citizens," Perspectives on Politics, vol. 12, issue 3, September 2014, p. 564 - 581. [Online] Available at: <https://doi.org/10.1017/S1537592714001595>
- [5] V. Alsina and J. L. Martí, "The Birth of the CrowdLaw Movement," Analyse & Kritik, vol. 40 no. 2, p. 337-358, 2018. [Online] Available at: https://www.analyse-und-kritik.net/Dateien/5be9b068e75b1_alsina_marti.pdf
- [6] B. S. Noveck, "Crowdlaw: Collective Intelligence and Lawmaking," Analyse & Kritik, vol. 40 no. 2, p. 359-380. [Online] Available at: https://www.analyse-und-kritik.net/Dateien/5be9b083bc696_noveck.pdf
- [7] B. S. Noveck et al. "CROWDLAW FOR CONGRESS: Strategies for 21st Century Lawmaking," The GovLab, October 2020. [Online] Available at: https://congress.crowd.law/files/crowdlaw_playbook_Oct2020.pdf
- [8] B.S. Noveck et al. "From Citizen to Senator: Artificial Intelligence and the Reinvention of Citizen Lawmaking in Brazil," The GovLab, April 2025. [Online] Available at: <https://files.thegovlab.org/from-citizen-to-senator.pdf>
- [9] J. Mathews, "The Little Democratic Engine That Could," Zocalo Public Square, April 15, 2025. [Online] Available at: <https://www.zocalopublicsquare.org/the-little-democratic-engine-that-could-engaged-california/>
- [10] M. Ryan, D. Gambrell, and B. S. Noveck, "Using Collective Intelligence to Solve Public Problems," Nesta & The GovLab, October 2020. [Online] Available at: https://directus.thegovlab.com/uploads/ci_cases/originals/b3b84db2-296c-435d-968a-8d066c6a7b87.pdf
- [11] C. Summerfield et al. "How will advanced AI systems impact democracy?" August 2024. [Online] Available at: <https://arxiv.org/abs/2409.06729>
- [12] R. Bjarnason, D. Gambrell, and J. Lanthier-Welch, "Using Artificial Intelligence to Accelerate Collective Intelligence: Policy Synth and Smarter Crowdsourcing," June 2024. [Online] Available at: <https://doi.org/10.48550/arXiv.2407.13960>
- [13] N. S. T. Chong, "Deep Research: Revolutionizing Information Synthesis with AI," United Nations University, February 20, 2025. [Online] Available at: <https://c3.unu.edu/blog/deep-research-revolutionizing-information-synthesis-with-ai>
- [14] M. J. Salganik and K. E. C. Levy, "Wiki Surveys: Open and Quantifiable Social Data Collection," PLoS ONE vol. 10 no. 5, May 2015. [Online] Available at: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0123483>
- [15] A. Dinesh, "Smarter crowdsourcing to tackle COVID-19: Beyond the open call," in The Routledge Handbook of Collective Intelligence for Democracy and Governance (1st Edition), Routledge, 2023. [Online] Available at: <https://www.taylorfrancis.com/chapters/oa-edit/10.4324/9781003215929-44/smarter-crowdsourcing-tackle-covid-19-anirudh-dinesh>
- [16] B. Goldberg et al. "AI and the Future of Digital Public Squares," December 2024. [Online] Available at: <https://arxiv.org/abs/2412.09988>
- [17] D. Gambrell. "AI and You: How New Jersey Is Using Collective Intelligence to Get Smarter on AI," Reboot Democracy, July 30, 2024. [Online] Available at: <https://rebootdemocracy.ai/blog/ai-and-you-all-our-ideas-new-jersey>
- [18] B.S. Noveck, "How AI could restore our faith in democracy," Fast Company, January 9, 2024. [Online] Available at: <https://www.fastcompany.com/91001497/ai-faith-in-democracy>
- [19] Report to the Governor on Artificial Intelligence," New Jersey AI Task Force, November 2024. [Online] Available at: <https://innovation.nj.gov/news/NJ-AI-Task-Force-Report.pdf>
- [20] S. Fox-Sowell, "New Jersey launches generative AI assistant and training tool for state employees," StateScoop, July 5, 2024. [Online] Available here: <https://statescoop.com/new-jersey-ai-assistant-training-tool-state-employees/>
- [21] S. Marcucci and S. Verhulst, "Reimagining the Policy Cycle in the Age of Artificial Intelligence," SSRN, February 2025. [Online] Available at: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5137557
- [22] A. Asatiani, P. Malo, P. R. Nagbøl, E. Penttinen, T. Rinta-Kahila, A. Salovaara "Challenges of Explaining the Behavior of Black-Box AI Systems," MIS Quarterly Executive: vol. 19: iss. 4, Article 7 (2020).. [Online] Available at: <https://aisel.aisnet.org/misqe/vol19/iss4/7>
- [23] C. Rudin, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," Nature Machine Intelligence vol 1, p. 206-215 (2019). [Online] Available at: <https://www.nature.com/articles/s42256-019-0048-x>
- [24] "OpenAI o3 and o4-mini System Card," OpenAI, April 16, 2025. [Online] Available at: <https://cdn.openai.com/pdf/2221c875-02dc-4789-800b-e7758f3722c1/o3-and-o4-mini-system-card.pdf>
- [25] Z. Schiffer, E. Mullin, W. Knight "OpenAI and the FDA Are Holding Talks About Using AI In Drug Evaluation," May 7, 2025. [Online] Available at: <https://www.wired.com/story/openai-fda-doge-ai-drug-evaluation/>
- [26] D. Kayyali, "Ask the Experts: AI Surveillance and US Immigration Enforcement," April 22, 2025. [Online] Available at: <https://www.techpolicy.press/ask-the-experts-ai-surveillance-and-us-immigration-enforcement/>
- [27] N. Jones, "OpenAI's 'deep research' tool: is it useful for scientists?" Nature, February 6, 2025. [Online] Available at: <https://doi-org.ezproxy.neu.edu/10.1038/d41586-025-00377-9>
- [28] Z. Luo, "The Rising Threat to Emerging AI-Powered Search Engines," February 7, 2025. [Online] Available at: <https://arxiv.org/abs/2502.04951>
- [29] K. Jazwińska and A. Chandrasekar, "AI Search Has A Citation Problem," Columbia Journalism Review, March 6, 2025. [Online] Available at: https://www.cjr.org/tow_center/we-compared-eight-ai-search-engines-theyre-all-bad-at-citing-news.php
- [30] D. Gilbert and V. Elliott, "DOGE put a college student in charge of using AI to rewrite regulations," WIRED, April 30, 2025. [Online] Available at: <https://www.wired.com/story/doge-college-student-ai-rewrite-regulations-deregulation/>
- [31] D. C. Youvan, "The Future of Knowledge: AI as a Living Publication and the Evolution of Research Dissemination," February 11, 2025. [Online] Available at: <http://dx.doi.org/10.13140/RG.2.2.26653.22244>
- [32] S. Longpre, "AI crawler wars threaten to make the web more closed for everyone," February 11, 2025. [Online] Available at: <https://www.technologyreview.com/2025/02/11/1111518/ai-crawler-wars-closed-web/>
- [33] "Impact Report 2024," New Jersey Office of Innovation, September 2024. [Online] Available at: <https://innovation.nj.gov/impact-report/2024/>
- [34] M. Ryan, D. Gambrell, and B. S. Noveck, "Collective Intelligence: A checklist for the public sector," Nesta & The GovLab, October 2020. [Online] Available at: https://media.nesta.org.uk/documents/Collective_Intelligence_A_checklist_for_the_public_sector.pdf
- [35] D. Berliner, "What AI Can't Do for Democracy," Boston Review, November 21, 2024. [Online] Available at: <https://www.bostonreview.net/articles/what-ai-cant-do-for-democracy/>

APPENDIX

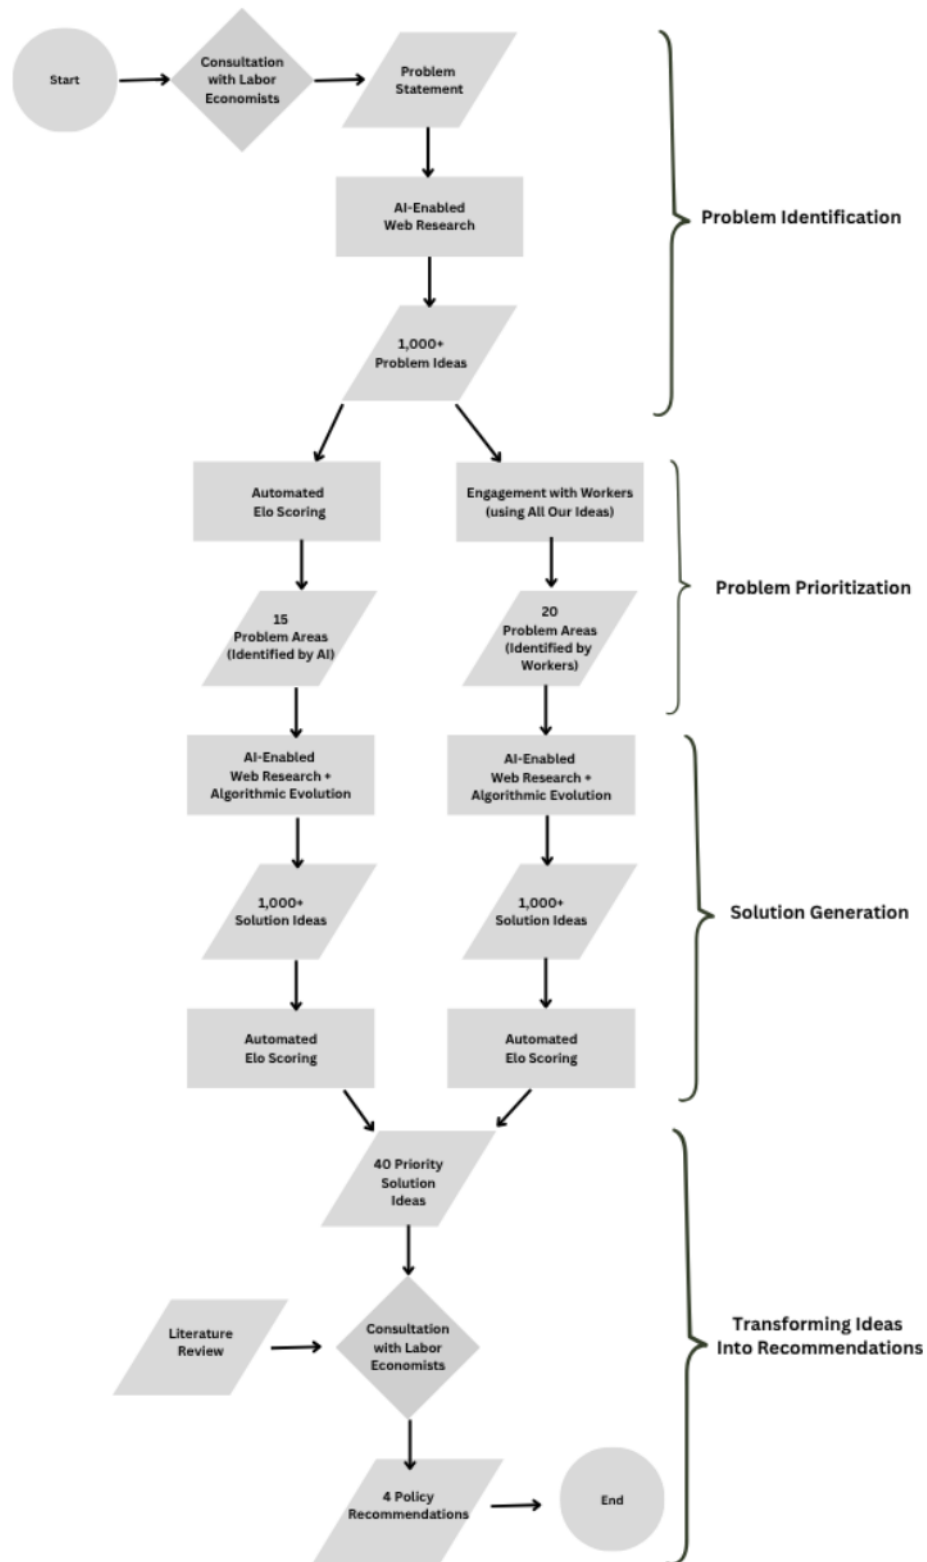


Fig. 3. Flowchart showing the inputs and outputs at each stage of the Working Group's policymaking process.